

Basic Neural Networks

OCTOBER 18, 2017 — WORKSHOP #1

Intro

We want to generate a prediction from some input

Output should give some probability of all our expected outputs

Network should give highest probability to correct class most of the time

Forward Pass

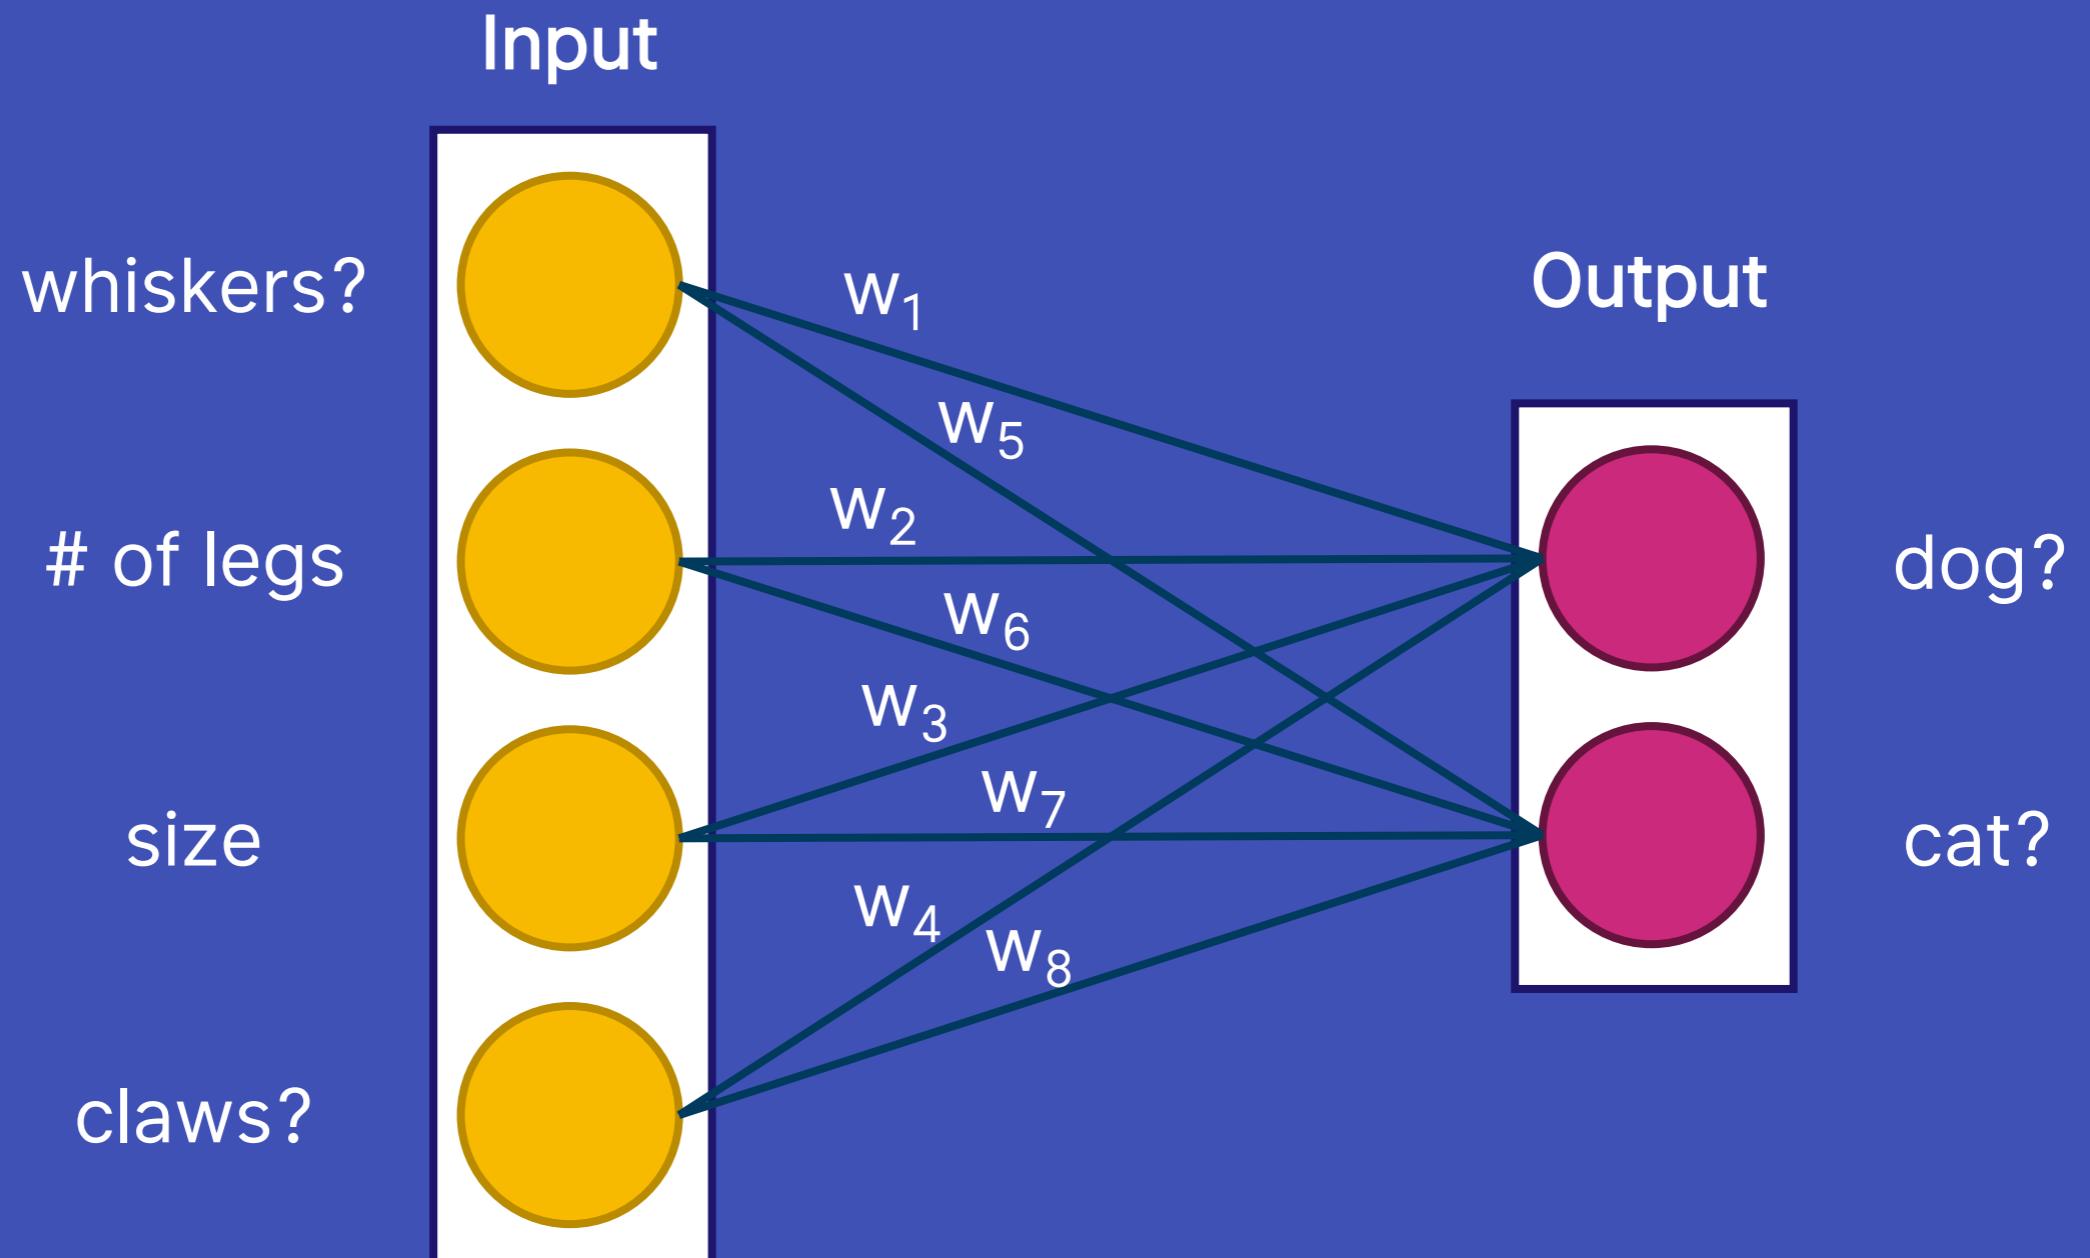
Fully-Connected Layers

Fully-connected layers are computations where every input neuron is connected to every output neuron

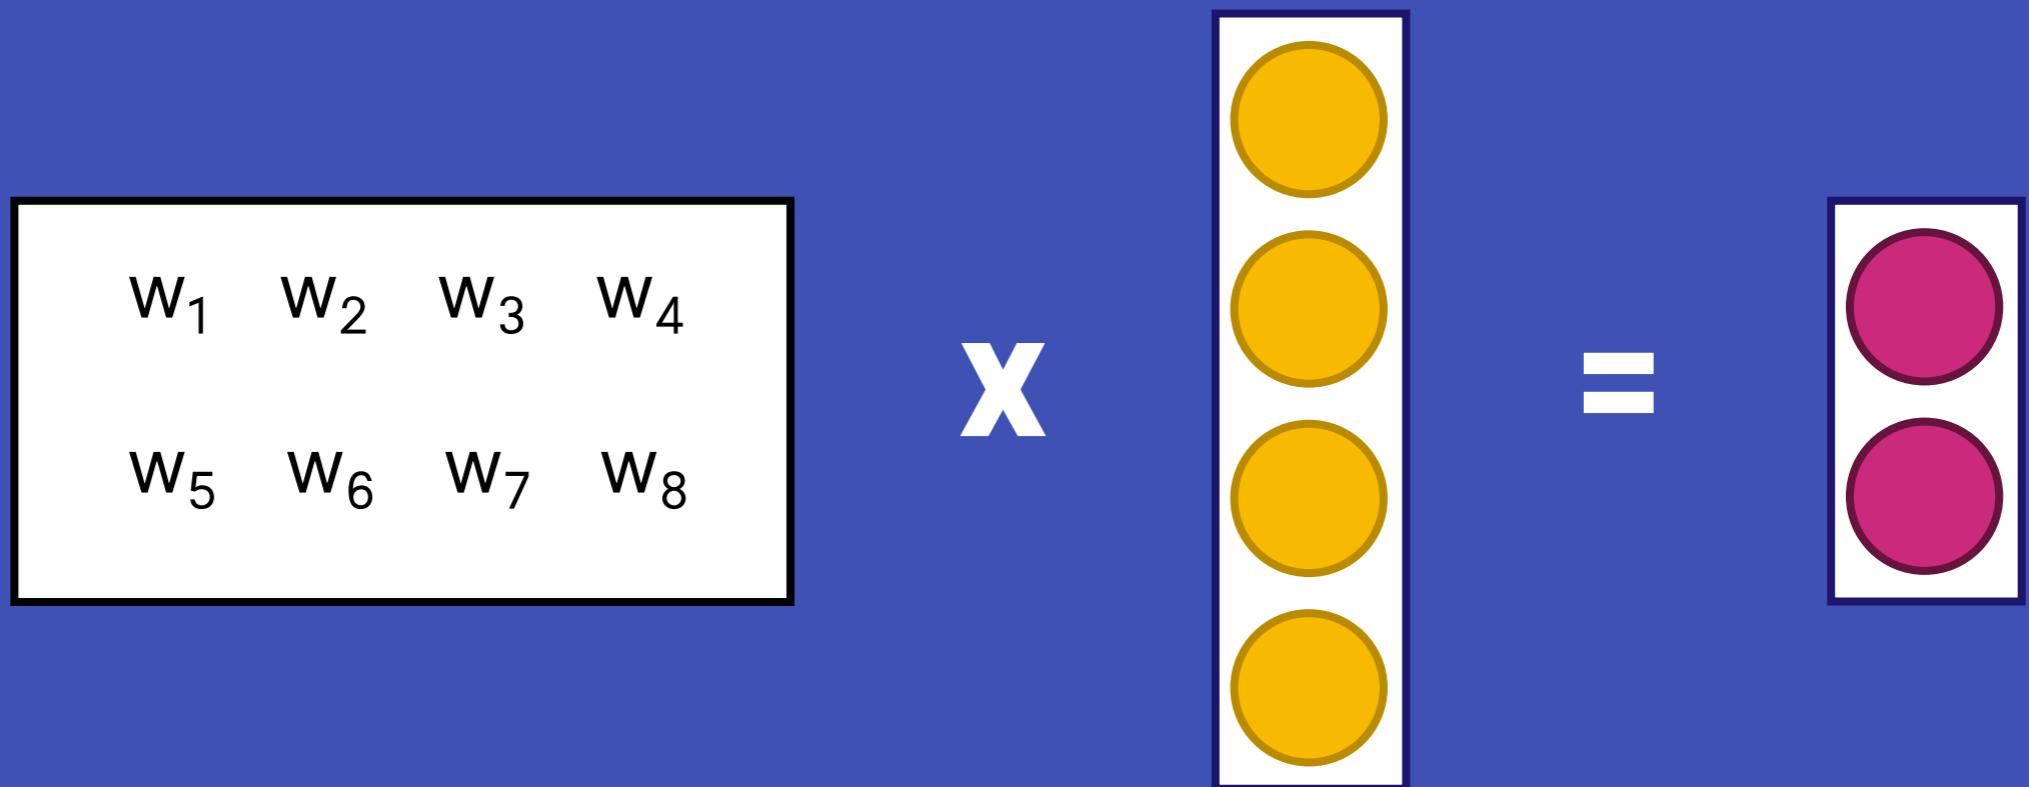
We transform the input by using weights to transform it into the output

The output will have a different shape than the input

Fully-Connected Network



FC Net as Matrix-Vec Mult



Probability Distribution

The output of our network is arbitrarily scaled

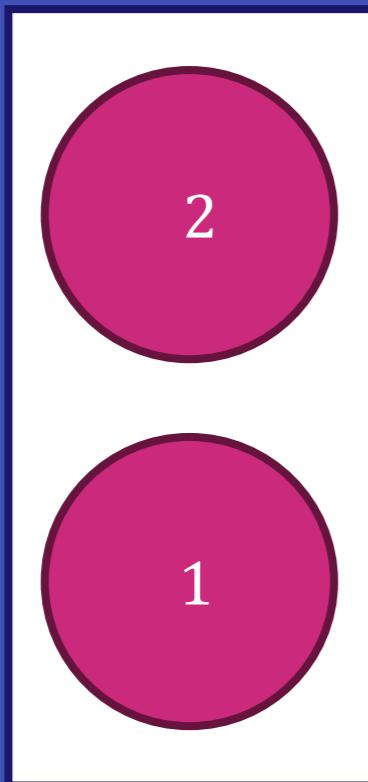
We want to turn the arbitrary scale into a proper probability distribution

We do so by using the softmax calculation:

$$p_i = \frac{e^{o_i}}{\sum_j e^{o_j}}$$

Softmax Layer

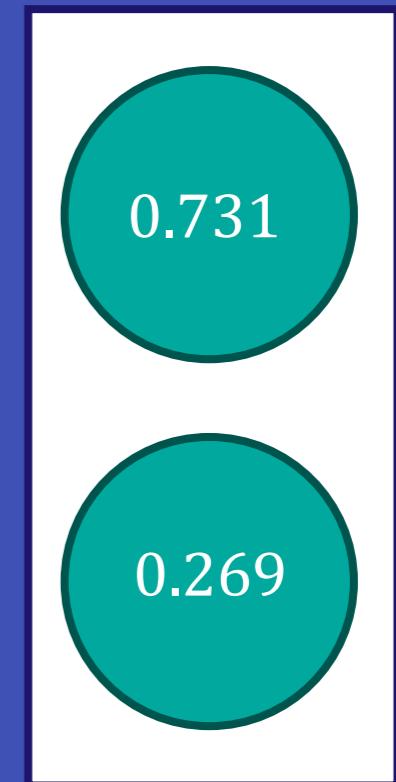
Output



Softmax

A white rectangular box containing two teal circles, each with a black outline. The top circle contains the formula $\frac{e^2}{e^2 + e^1}$ and the bottom circle contains the formula $\frac{e^1}{e^2 + e^1}$.

=



Loss Value

The network needs an “objective” to work towards—being accurate

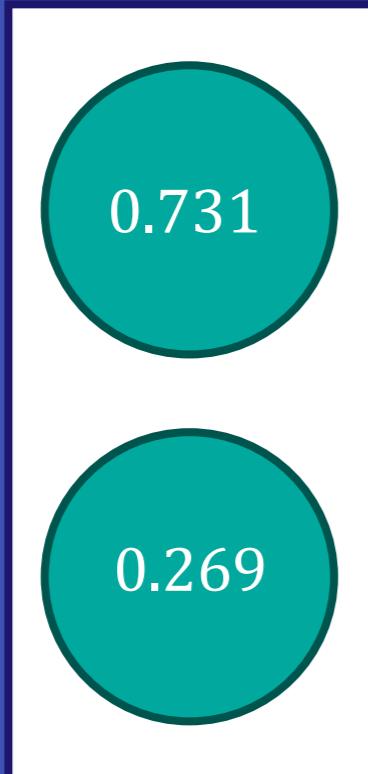
We use the targets to figure out how “incorrect” the network was

The loss metric we use is cross-entropy loss:

$$L = -\log(pc)$$

Softmax Layer

Softmax



Loss

$$L = -\log(0.731) = 0.313$$

Data Splitting

We are interested in how our model does on unseen data above all

We can split data into training and test—but if we optimize on test, test becomes “indirectly” seen

Instead, we split into train, validation, and test; we run test only once

EXERCISE

Code Up the Forward Pass

Split the dataset into train, validation, and test

Initialize the parameters to random

Create the linear and softmax cross-entropy layers

Backward Pass

Backpropagation

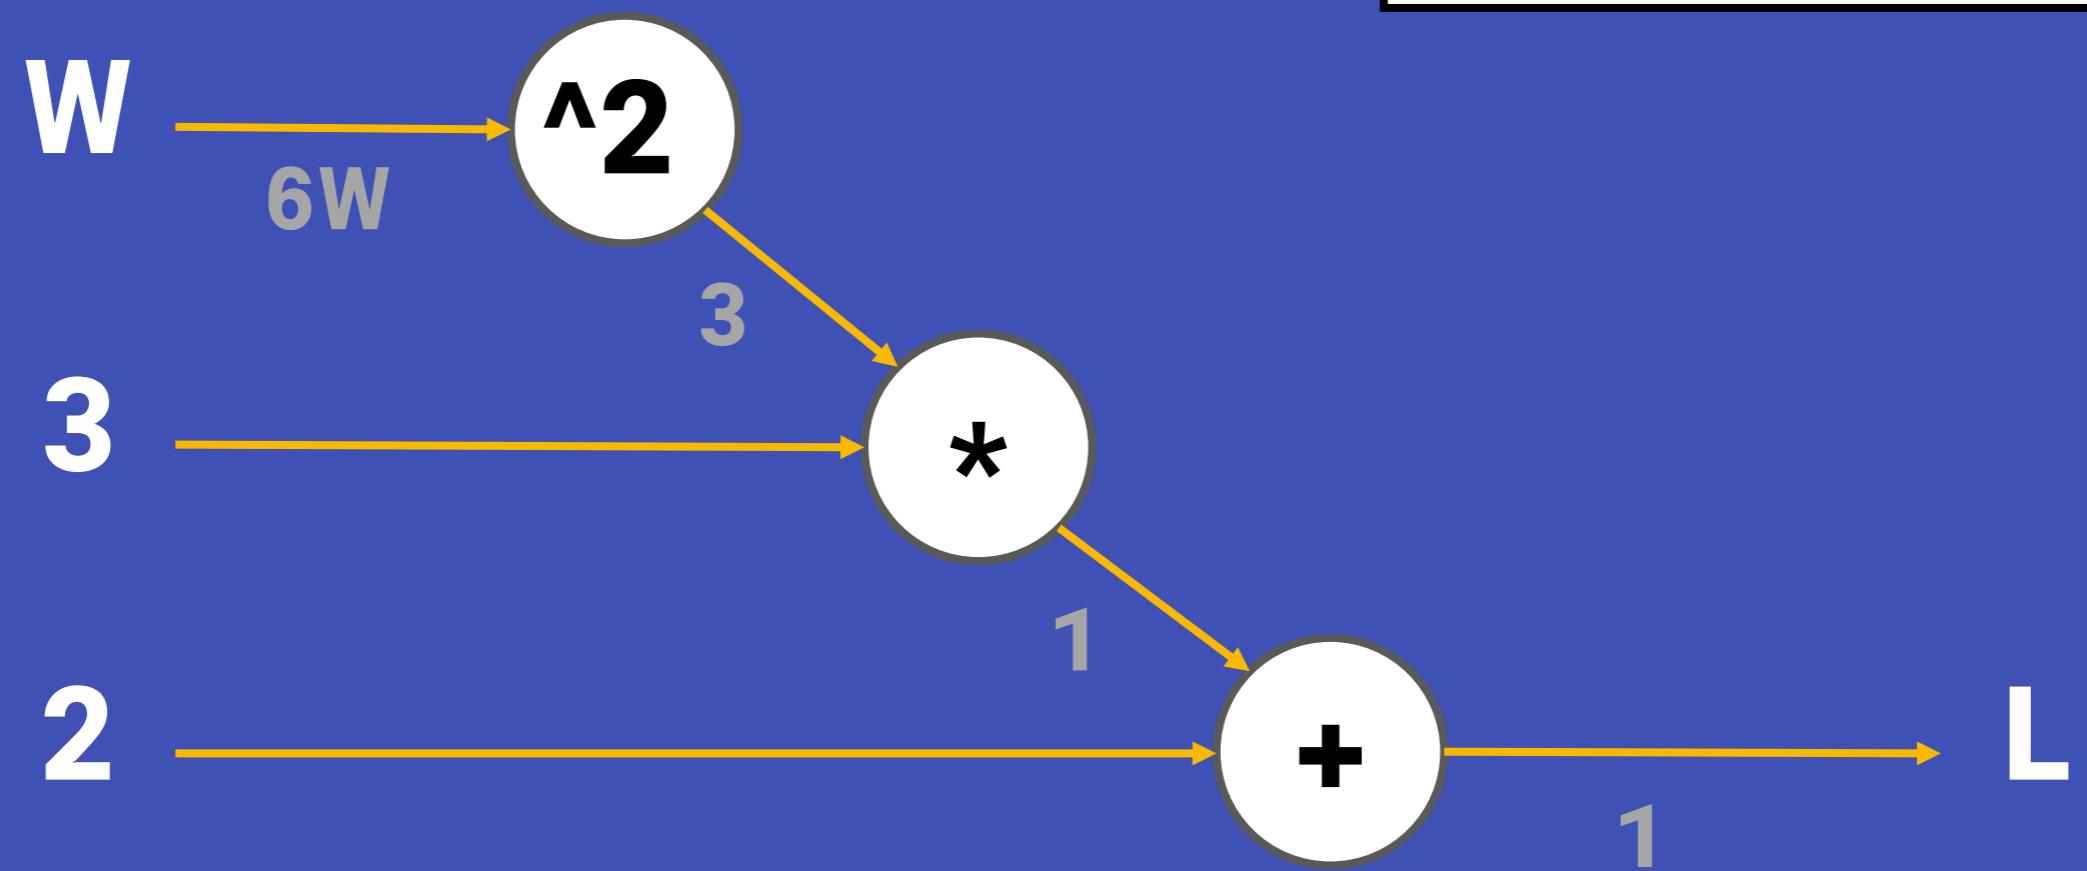
Each computation is represented as a node in a computational graph

Gradients are computed at each node of the graph

Chain rule is applied recursively to get total gradients on each weight

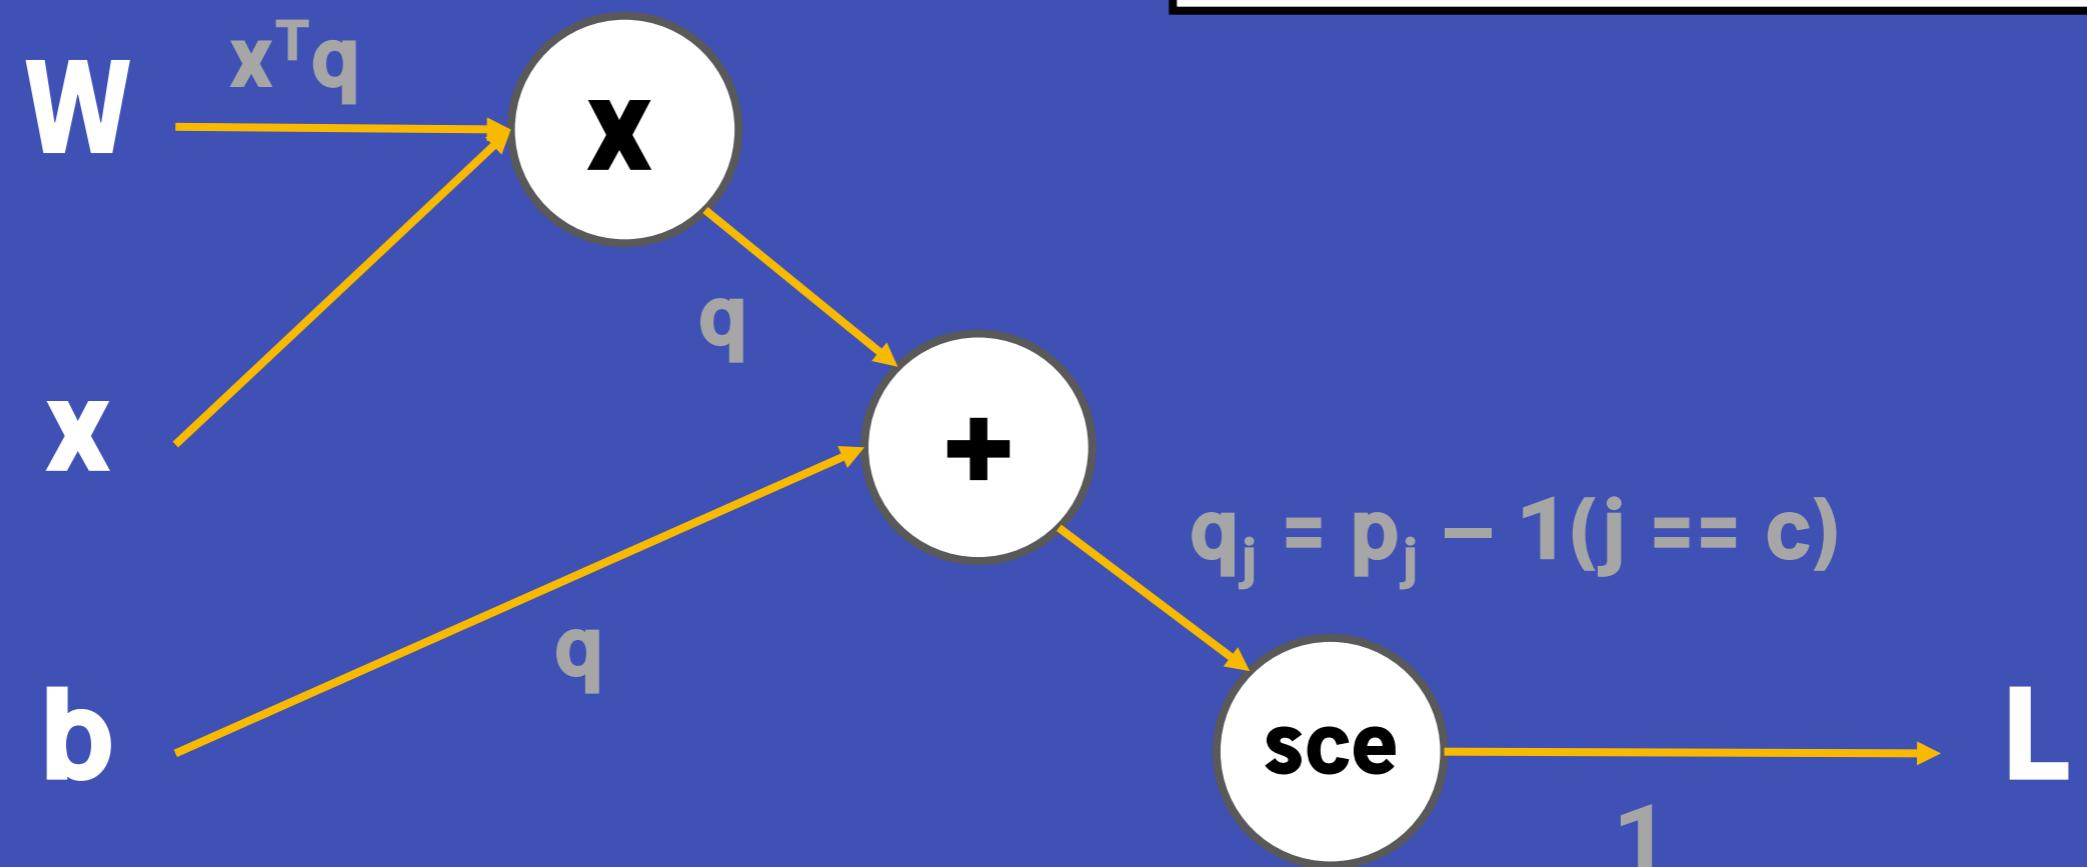
Simple Backprop

$$L = 3W^2 + 2$$



Multidimensional Backprop

$$L = \text{sce}(Wx + b)$$



EXERCISE

Code Up the Backward Pass

Create the backward pass through the linear layer

Create the backward pass through the softmax cross-entropy layer

Training

Weight Update

Once we find gradients, we must change each weight in that direction

We can use a update rules to accomplish this; the simplest is SGD:

$$W_i := W_i - \alpha \left(\frac{\partial L}{\partial W_i} \right)$$

Training Iterations

Each time we move the weights in the direction of the negative gradient, we get closer to a good value for the weights

We want to do this many times until we arrive at good weight values (i.e. converge)

We also want to periodically check the network's accuracy on the training and validation sets

EXERCISE

Code Up the Training Regime

Code the training iteration loop

Implement SGD weight updates

Periodically log the training and validation losses

Conclusion

OCTOBER 26

Animesh Garg

Deep reinforcement learning
with computer vision for robotics

Applications include surgical
robotics, robotic navigation, and
more!



NOVEMBER 9

Bharath Ramsundar

Deep learning for healthcare

Applications include drug discovery,
chemical retrosynthesis, and more!



Thank You!

Please fill out our survey!

See you next week